

NAME : G RAJU HT NO. : 2403A52018

Lab 8: L1 Regularization -Lasso: Implementing Lasso Regression using libraries and comparing results with Linear Regression. Kaggle Dataset Link: <https://www.kaggle.com/datasets/debjayotipodder/co2-emission-by-vehicles> Tasks:

1. Load dataset and select features: Engine Size, Cylinders, Fuel Consumption Target: CO2 Emissions
2. Apply StandardScaler.
3. Train Linear Regression and Lasso Regression.
4. Use Ridge with multiple alpha values and find best alpha using: GridSearchCV
5. Report best model performance and compare with Linear Regression.

```
import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.metrics import r2_score, mean_squared_error

import warnings
warnings.filterwarnings("ignore")

print("Libraries Imported Successfully")
```

Libraries Imported Successfully

```
# Load dataset
file_path = "/content/CO2_Emissions_Canada.csv"
df = pd.read_csv(file_path)

print("Dataset Shape:", df.shape)
print("\nColumns:\n", df.columns)
df.head()
```

Dataset Shape: (7385, 12)

Columns:

```
Index(['Make', 'Model', 'Vehicle Class', 'Engine Size(L)', 'Cylinders',
       'Transmission', 'Fuel Type', 'Fuel Consumption City (L/100 km)',
       'Fuel Consumption Hwy (L/100 km)', 'Fuel Consumption Comb (L/100 km)',
       'Fuel Consumption Comb (mpg)', 'CO2 Emissions(g/km)'],
      dtype='object')
```

	Make	Model	Vehicle Class	Engine Size(L)	Cylinders	Transmission	Fuel Type	Fuel Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	Fuel Consumption Comb (mpg)	Emissi
0	ACURA	ILX	COMPACT	2.0	4	AS5	Z	9.9	6.7	8.5	33	
1	ACURA	ILX	COMPACT	2.4	4	M6	Z	11.2	7.7	9.6	29	
2	ACURA	ILX	HYBRID	1.5	4	AV7	Z	6.0	5.8	5.9	48	
3	ACURA	MDX	SUV - SMALL	3.5	6	AS6	Z	12.7	9.1	11.1	25	
	PDY	Q5										

```
# Selecting required features
features = ["Engine Size(L)",
            "Cylinders",
            "Fuel Consumption Comb (L/100 km)"]
```

target = "CO2 Emissions(g/km)"

```
X = df[features]
y = df[target]
```

```
print("Feature Matrix Shape:", X.shape)
print("Target Shape:", y.shape)
```

Feature Matrix Shape: (7385, 3)  
Target Shape: (7385,)

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

print("Training Data Shape:", X_train.shape)
print("Testing Data Shape:", X_test.shape)

Training Data Shape: (5908, 3)
Testing Data Shape: (1477, 3)
```

```
scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Standard Scaling Applied Successfully")

Standard Scaling Applied Successfully
```

```
lr = LinearRegression()
lr.fit(X_train_scaled, y_train)

y_pred_lr = lr.predict(X_test_scaled)

r2_lr = r2_score(y_test, y_pred_lr)
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))

print("Linear Regression Results")
print("R2 Score :", round(r2_lr, 4))
print("RMSE      :", round(rmse_lr, 4))

Linear Regression Results
R2 Score : 0.8773
RMSE      : 20.5407
```

```
lasso = Lasso(alpha=0.01) # Small alpha value
lasso.fit(X_train_scaled, y_train)

y_pred_lasso = lasso.predict(X_test_scaled)

r2_lasso = r2_score(y_test, y_pred_lasso)
rmse_lasso = np.sqrt(mean_squared_error(y_test, y_pred_lasso))

print("Lasso Regression Results")
print("R2 Score :", round(r2_lasso, 4))
print("RMSE      :", round(rmse_lasso, 4))

Lasso Regression Results
R2 Score : 0.8773
RMSE      : 20.5409
```

```
ridge = Ridge()

param_grid = {
    "alpha": [0.001, 0.01, 0.1, 1, 10, 100, 500, 1000]
}

grid = GridSearchCV(ridge, param_grid, cv=5, scoring="r2")
grid.fit(X_train_scaled, y_train)

best_alpha = grid.best_params_["alpha"]
best_ridge = grid.best_estimator_

y_pred_ridge = best_ridge.predict(X_test_scaled)

r2_ridge = r2_score(y_test, y_pred_ridge)
rmse_ridge = np.sqrt(mean_squared_error(y_test, y_pred_ridge))

print("Ridge Regression Results")
print("Best Alpha :", best_alpha)
print("R2 Score   :", round(r2_ridge, 4))
print("RMSE       :", round(rmse_ridge, 4))

Ridge Regression Results
Best Alpha : 10
```

```
R2 Score : 0.8773
RMSE : 20.5414
```

```
print("\n===== MODEL COMPARISON =====")

print("Linear Regression -> R2:", round(r2_lr,4), " | RMSE:", round(rmse_lr,4))
print("Lasso Regression -> R2:", round(r2_lasso,4), " | RMSE:", round(rmse_lasso,4))
print("Ridge Regression -> R2:", round(r2_ridge,4), " | RMSE:", round(rmse_ridge,4))

best_model = max(
    [("Linear Regression", r2_lr),
     ("Lasso Regression", r2_lasso),
     ("Ridge Regression", r2_ridge)],
    key=lambda x: x[1]
)

print("\nBest Model Based on R2 Score:", best_model[0])
```

```
===== MODEL COMPARISON =====
Linear Regression -> R2: 0.8773 | RMSE: 20.5407
Lasso Regression -> R2: 0.8773 | RMSE: 20.5409
Ridge Regression -> R2: 0.8773 | RMSE: 20.5414
```

Best Model Based on R2 Score: Linear Regression

```
import matplotlib.pyplot as plt
import numpy as np

plt.figure()

# Scatter plot of actual vs predicted
plt.scatter(y_test, y_pred_lr)

# Perfect prediction line (diagonal line)
min_val = min(y_test.min(), y_pred_lr.min())
max_val = max(y_test.max(), y_pred_lr.max())
plt.plot([min_val, max_val], [min_val, max_val])

plt.xlabel("Actual CO2 Emissions (g/km)")
plt.ylabel("Predicted CO2 Emissions (g/km)")
plt.title("Actual vs Predicted CO2 Emissions")

plt.show()
```

